

1 Design with Azure

Main components:

- Azure functions
- Azure durable functions
- Azure blob storage, Azure cache for Redis

Main problems encountered:

- RAM limitations in each function
- Intermediate aggregations of data structures after each phase
- Azure functions can be written in Rust using custom handlers, but Azure durable functions cannot be written using Rust. (We are not 100% sure but, by looking on the internet we did not find any official information about writing durable functions using Rust in Azure)
- It seems to be impossible to use an orchestrator for example written in python to orchestrate functions written in Rust since we are in the same function app, thus same runtime ([See link](#))

Possible solutions:

- Use only Azure functions written in Rust and orchestrate them without using durable functions but exploiting other Azure services such as message queues to re-create the fan-in fan-out pattern we want to have. (We haven't yet dug deep in this solution)
- Write Azure functions in Rust using custom handlers and trigger them using HTTP requests. We can run a client orchestrator on a docker container hosted in the clouds that act as a stateful function and orchestrate the workload by triggering our functions with HTTP requests.
- Switch to Java

2 Azure serverless design for parallelizing the algorithm

2.1 Main issues in parallelizing

- Passing data between functions that are in different phases of the algorithm:
 - Since data structures may be quite big, we can't simply exploit the single orchestrator to handle all the different partitions and pass them as parameters to functions of different phases.¹
To solve this, we thought about using **Azure cache for redis** (in-memory key-value noSQL database) as storage medium to pass data through different phases.
- Why entity-based partitions and not simple file partitioning ?
 - Assuming we have to scale to big graphs
 - Mainly because this allows workers to work on local versions of data structures. If we simply partitioned the file in N split without any semantics, then we would have had to keep all the structures global, thus use something like a queryable DB to share them (read-only) with all our functions (since a single function cannot host in memory the full Ψ_{ETD} data structure for example)
 - Still, we may have to deal with some memory problems in the functions where we merge the data structures. However, in our algorithm, we will never have to merge the full Ψ_{ETD} data structure with also the full information about Ψ_{ETPD} , which we think is the most memory critical one.

2.2 Phase 0 - How to partition the starting file

Generate N **entity based partitions**, i.e. each of our parallel worker (function) will be assigned a file containing only triplets where the **subject** is part of the set of entities assigned to the partition (e.g. Worker 1 will be assigned with set of entities {Alice, Luca, Michele}, thus the partition will contain only triplets with Alice, Luca, and Michele as **subject**). We generate these **entity based partitions** in such a way that each worker works roughly on the same amount of data.

2.3 Phase 1

Each worker will perform computations on its assigned split to generate local data structures Ψ_{ETD}, Ψ_{CEC} . After phase 1, however, we are forced to merge all the different local Ψ_{ETD} data structures, because information about the mapping **Entity** \rightarrow **Set of types** is needed during phase 2 for the entity constraints extraction. Anyway, this should not be too much memory consuming, since in Ψ_{ETD} we still don't have any data in the Ψ_{ETPD} map in each entity (this will be populated locally during phase 2).

¹We haven't implemented anything yet, using our common sense we simply thought that passing \approx GB data structures as function parameters was not a good idea

2.4 Phase 2

Up to now we have N different local Ψ_{CEC} (one for each worker), and Ψ_{ETD} which is the same for each worker. After this phase we will have again different local versions of Ψ_{ETD} (with the additional local informations about entity constraints (Ψ_{ETPD})), Ψ_{CEC} , and Ψ_{CTP}

- Do it as on paper for computing a local version of Ψ_{ETD}

2.5 Phase 3

Now each worker will use its local Ψ_{ETD} to compute local versions of Ψ_{SUPP} and Ψ_{PTT} . We still cannot compute Ψ_{CONF} during this phase as in the algorithm described in the paper, since we still have to merge together all the local versions of Ψ_{CEC} .

After phase 3 we now need to merge together the local versions of Ψ_{SUPP} , Ψ_{PTT} , Ψ_{CEC} , Ψ_{CTP} and use the merged Ψ_{SUPP} and Ψ_{CEC} to compute Ψ_{CONF} . These 3 data structures (SUPP, PTT, CEC) should not be memory critical anymore.

2.6 Phase 4

During phase 4 we compute Ψ_{CONF} using Ψ_{SUPP} and Ψ_{CEC} . We don't have memory issues since there is no need to keep the full Ψ_{ETD} data structure in memory, which was the most critic one.

Then we perform shapes extraction as in the QSE-exact algorithm using Ψ_{CTP} :

- Easily parallelizable by running functions over a subsets of triplets that check if the triplet is higher than ϵ or ω based if we are running on confidence or support

2.7 Graphical explanation

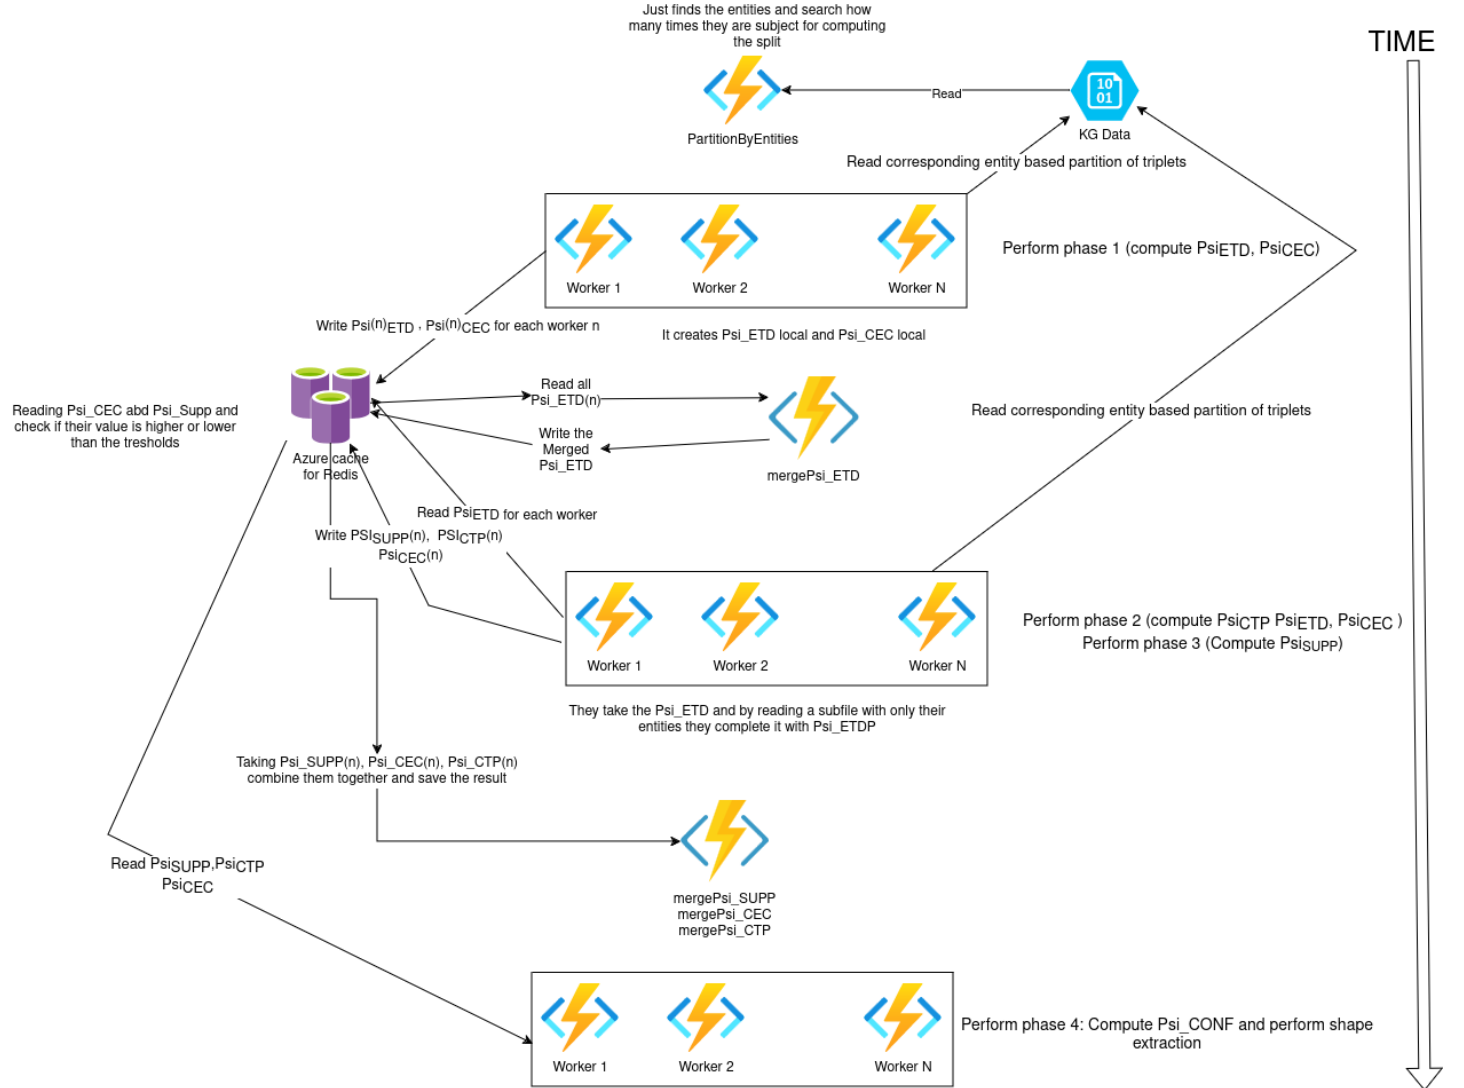


Figure 1: A sample image

3 LINKS

3.1 PAPER

- Original paper
- Summary of the algorithms

3.2 WHY RUST

<https://blog.scanner.dev/serverless-speed-rust-vs-go-java-python-in-aws-lambda-functions/>

3.3 AWS

- AWS RUST
- AWS itembatcher
- AWS dynamic parallelism used
- AWS parallel

3.4 AZURE

- Durable Functions
- HOW TO create durable functions Java with VS code